

南京大学本科生实验报告

课程名称：计算机网络

任课教师：田臣/李文中

助教：

| | | | |
|-------|--|---------|--------------------|
| 学院 | 计算机科学与技术系 | 专业（方向） | 计算机科学与技术 |
| 学号 | 202220013 | 姓名 | 徐简 |
| Email | 161200063@smail.nju.edu.cn | 开始/完成日期 | 2021.5.8-2021.5.19 |

1. 实验名称

Lab 5: Respond to ICMP

2. 实验目的

1. Respond to ICMP messages like echo requests ("pings").
2. Generate ICMP error messages

3. 实验内容、代码与结果

Step 1: Initialize required tables

- Define required classes.
- Initialize a `forward_table` and a `packet_queue`

`ftItem` is exactly the same as what I used in lab4, while there is a slight change in `queueItem`. To generate ICMP error messages, the `ifaceName` of the packet needs to be recorded. Therefore, add an `icmp_info` attribute to `queueItem`, indicating the coming port of the packet.

```
class queueItem():
    def __init__(self, pkt, match, icmp):
        self.pkt=pkt
        self.rounds=0
        self.time=0
        self.match=match
        self.icmp_info=icmp # icmp_info records the coming port of the packet
```

Step 2: Responding to ICMP echo requests

Define a function to respond to ICMP echo requests. The function takes in the original request packet and the IP address of the destination, and returns a reply packet.

The details of building each header can be found in the switchyard API references.

```
def ping_reply(packet, dstip):
    ether=Ethernet()
    ether.ethertype=EtherType.IP
    # construct ether header
    ip=IPv4()
    ip.src=IPAddr(dstip)# the reply ip src is ping dst
    ip.dst=IPAddr(packet[IPv4].src)# the reply ip dst is ping src
    ip.protocol=IPProtocol.ICMP
    ip.ttl=64
    ip.ipid=0
    #construct ip header
    icmp=ICMP()
    icmp.icmptype=ICMPType.EchoReply
    icmp.icmpcode=ICMPCodeEchoReply.EchoReply
    icmp.icmpdata.sequence=packet[ICMP].icmpdata.sequence
    icmp.icmpdata.identifier=packet[ICMP].icmpdata.identifier
    icmp.icmpdata.data=packet[ICMP].icmpdata.data
    #construct icmp header
    return ether+ip+icmp
```

Then, we need to call the function correctly when the router receives packets. There is only one case when the router needs to generate an ICMP echo reply, which is that the router receives a IPv4 packet heading to one of its interfaces and the packet happens to be an ICMP echo request. In other cases, the function above will not be called.

```
if ipv4:
    head=packet[IPv4]
    for i in self.ip_list:
        if packet[IPv4].dst==i:# IPv4 packet heading to the router interfaces
            if packet.has_header(ICMP) and
packet[ICMP].icmptype==ICMPType.EchoRequest:
                # the packet happens to be a ICMP echo request
                packet=ping_reply(packet,i)
                head=packet[IPv4]
                break
```

Above are the logic of responding the ICMP echo requests.

Step 3: Generating ICMP error messages

Just like Step3, define a function to generate ICMP error messages. The function needs the IP source and IP destination of the message, icmp type, icmp code, ttl and the packet that causes an ICMP error messages.

```
def construct_icmperror(ipsrc,ipdst,xtype,xcode=0,ttl=64,origpkt=None):
    icmp = ICMP()
    icmp.icmptype = xtype
    icmp.icmpcode=xcode
    if not origpkt is None:
        i = origpkt.get_header_index(Ethernet)
        del origpkt[i]
        icmp.icmpdata.data = origpkt.to_bytes()[28:]
        icmp.icmpdata.origdgramlen = len(origpkt)
    #set icmptype, icmpcode and fill icmpdata with the original packet
    ip = IPv4()
    ip.protocol = IPPROTO.ICMP
    ip.src=IPAddr(ipsrc)
    ip.dst=IPAddr(ipdst)
    ip.ttl=ttl
    # set IPv4 header
    ether=Ethernet()
    #set Ethernet header
    pkt = ether+ip + icmp
    #return the error message
    return pkt
```

According to the lab manual, there are 4 cases when we should generate ICMP error messages.

- An incoming packet is destined to an IP addresses assigned to one of the router's interfaces, but the packet is not an ICMP echo request

This case is the `else` side of sending an ICMP echo reply.

```
if packet[IPv4].dst==i:
    if packet.has_header(ICMP) and packet[ICMP].icmptype==ICMPType.EchoRequest:
        ...
    ping reply
    ...
else:
    # receive a packet heading to the router itself, but not a ping request
    for i in self.interfaces:
        if i.name==ifaceName:
            p=i
            break
    #p is the dst interface of the router
```

```

        #ip address of interface p is the src, ip address of the packet is the
dst
        # icmp type destination port unreachable

packet=construct_icmperror(p.ipaddr,head.src,ICMPType.DestinationUnreachable,3,
64,packet)

```

- When attempting to match the destination address of an IP packet with entries in the forwarding table, no matching entries are found

```

# cannot match in the forwarding table
if pos == -1:
    print("cannot match?")
    for i in self.interfaces:
        if i.name==ifaceName:
            p2=i
            break
packet=construct_icmperror(p2.ipaddr,head.src,ICMPType.DestinationUnreachable,0
,64,packet)
    # generate a error message
    #look up the forwarding table for the error message packet
    head=packet[IPv4]
    pos=-1
    maxprifixlen=-1
    index=0
    for i in self.forward_table:
        if((int(head.dst)&int(i.mask))==int(i.prefix)):
            netaddr=IPv4Network(str(i.prefix)+"/"+str(i.mask))
            if netaddr.prefixlen>maxprifixlen:
                maxprifixlen=netaddr.prefixlen
                pos=index
            index+=1
    #find the match in the table and add it to the packet queue
    self.q.append(queueItem(packet,self.forward_table[pos],ifaceName))

```

- After decrementing an IP packet's TTL value as part of the forwarding process, the TTL becomes zero.

```

if head.ttl<=0:
    for i in self.interfaces:
        if i.name==ifaceName:
            p1=i
            break# find the coming port
    packet=construct_icmperror(p1.ipaddr,head.src,ICMPType.TimeExceeded,0,64,pack
et)
    # then do maxprifix match for the error message packet

```

- ARP Failure.

```

if self.q[0].rounds>=5:
    for i in self.interfaces:
        if i.name == self.q[0].icmp_info:
            p4=i
            break
    packet=construct_icmperror(p4.ipaddr,self.q[0].pkt[IPv4].src,ICMPType.DestinationUnreachable,1,64,self.q[0].pkt)
    # arp failure icmp error message
    head=packet[IPv4]
    ...

    maxprix match for the error message
    ...

    del(self.q[0])
    newq=queueItem(packet,self.forward_table[pos],p4.name)
    self.q.append(newq)

```

Above are the logic of Generating ICMP error messages.

Step 4: Test

```
$ swyard -t myrouter3_testscenario.srpy myrouter.py
```

Results for test scenario IP forwarding and ARP requester tests: 28 passed, 0 failed, 0 pending

Passed:

- 1 ICMP echo request (PING) for the router IP address 192.168.1.1 should arrive on router-eth0. This PING is directed at the router, and the router should respond with an ICMP echo reply.
- 2 Router should send an ARP request for 10.10.1.254 out router-eth1.
- 3 Router should receive ARP reply for 10.10.1.254 on router-eth1.
- 4 Router should send ICMP echo reply (PING) to 172.16.111.222 out router-eth1 (that's right: ping reply goes out a different interface than the request).
- 5 ICMP echo request (PING) for the router IP address 10.10.0.1 should arrive on router-eth1.
- 6 Router should send ICMP echo reply (PING) to 172.16.111.222 out router-eth1.
- 7 ICMP echo request (PING) for 10.100.1.1 with a TTL of 1 should arrive on router-eth1. The router should decrement

Step 5: Deploy

- `client# ping -c 1 -t 1 192.168.200.1`

Wireshark capture from client-eth0. The display filter is empty. The packet list shows four packets: 1. ARP request from 30.00:00:00:00:01 to Broadcast, 42 bytes. 2. ARP request from 40.00:00:00:00:03 to 30.00:00:00:00:01, 42 bytes. 3. ICMP Echo (ping) request from 10.1.1.1 to 192.168.200.1, 98 bytes. 4. ICMP Echo (ping) request from 10.1.1.2 to 10.1.1.1, 70 bytes. The packet details for packet 4 show 'Time-to-live exceeded (Time to live exceeded in transit)'.

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|-------------|-------------------|-------------------|----------|--------|--|
| 1 | 0.000000000 | 30:00:00:00:00:01 | Broadcast | ARP | 42 | Who has 10.1.1.2? Tell 10.1.1.1 |
| 2 | 0.087460544 | 40:00:00:00:00:03 | 30:00:00:00:00:01 | ARP | 42 | 10.1.1.2 is at 40:00:00:00:00:03 |
| 3 | 0.087471539 | 10.1.1.1 | 192.168.200.1 | ICMP | 98 | Echo (ping) request id=0x6248, seq=1/256, ttl=1 (no response found!) |
| 4 | 0.189261252 | 10.1.1.2 | 10.1.1.1 | ICMP | 70 | Time-to-live exceeded (Time to live exceeded in transit) |

The initial TTL in the ICMP packets is set to be 1, so an ICMP time exceeded error.

- `client# ping -c 1 172.16.1.1`

Wireshark capture from client-eth0. The display filter is empty. The packet list shows eight packets: 1. ARP request from 30.00:00:00:00:01 to Broadcast, 42 bytes. 2. ARP request from 40.00:00:00:00:03 to 30.00:00:00:00:01, 42 bytes. 3. ICMP Echo (ping) request from 10.1.1.1 to 192.168.200.1, 98 bytes. 4. ICMP Echo (ping) request from 10.1.1.2 to 10.1.1.1, 70 bytes. 5. ICMP Echo (ping) request from 10.1.1.1 to 172.16.1.1, 98 bytes. 6. ICMP Echo (ping) request from 10.1.1.2 to 10.1.1.1, 70 bytes. 7. ARP request from 30.00:00:00:00:01 to 40.00:00:00:00:03, 42 bytes. 8. ARP request from 40.00:00:00:00:03 to 30.00:00:00:00:01, 42 bytes. The packet details for packet 6 show 'Destination unreachable (Network unreachable)'.

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|--------------|-------------------|-------------------|----------|--------|---|
| 1 | 0.000000000 | 30:00:00:00:00:01 | Broadcast | ARP | 42 | Who has 10.1.1.2? Tell 10.1.1.1 |
| 2 | 0.087460544 | 40:00:00:00:00:03 | 30:00:00:00:00:01 | ARP | 42 | 10.1.1.2 is at 40:00:00:00:00:03 |
| 3 | 0.087471539 | 10.1.1.1 | 192.168.200.1 | ICMP | 98 | Echo (ping) request id=0x6248, seq=1/256, ttl=1 (no response found!) |
| 4 | 0.189261252 | 10.1.1.2 | 10.1.1.1 | ICMP | 70 | Time-to-live exceeded (Time to live exceeded in transit) |
| 5 | 89.595436677 | 10.1.1.1 | 172.16.1.1 | ICMP | 98 | Echo (ping) request id=0x6250, seq=1/256, ttl=64 (no response found!) |
| 6 | 89.658615355 | 10.1.1.2 | 10.1.1.1 | ICMP | 70 | Destination unreachable (Network unreachable) |
| 7 | 94.834081847 | 30:00:00:00:00:01 | 40:00:00:00:00:03 | ARP | 42 | Who has 10.1.1.2? Tell 10.1.1.1 |
| 8 | 94.924434738 | 40:00:00:00:00:03 | 30:00:00:00:00:01 | ARP | 42 | 10.1.1.2 is at 40:00:00:00:00:03 |

Wireshark capture from router-eth2. The display filter is empty. The packet list shows eight packets: 1. ARP request from 30.00:00:00:00:01 to Broadcast, 42 bytes. 2. ARP request from 40.00:00:00:00:03 to 30.00:00:00:00:01, 42 bytes. 3. ICMP Echo (ping) request from 10.1.1.1 to 192.168.200.1, 98 bytes. 4. ICMP Echo (ping) request from 10.1.1.2 to 10.1.1.1, 70 bytes. 5. ICMP Echo (ping) request from 10.1.1.1 to 172.16.1.1, 98 bytes. 6. ICMP Echo (ping) request from 10.1.1.2 to 10.1.1.1, 70 bytes. 7. ARP request from 30.00:00:00:00:01 to 40.00:00:00:00:03, 42 bytes. 8. ARP request from 40.00:00:00:00:03 to 30.00:00:00:00:01, 42 bytes. The packet details for packet 6 show 'Destination unreachable (Network unreachable)'.

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|--------------|-------------------|-------------------|----------|--------|---|
| 1 | 0.000000000 | 30:00:00:00:00:01 | Broadcast | ARP | 42 | Who has 10.1.1.2? Tell 10.1.1.1 |
| 2 | 0.087460544 | 40:00:00:00:00:03 | 30:00:00:00:00:01 | ARP | 42 | 10.1.1.2 is at 40:00:00:00:00:03 |
| 3 | 0.087462548 | 10.1.1.1 | 192.168.200.1 | ICMP | 98 | Echo (ping) request id=0x6248, seq=1/256, ttl=1 (no response found!) |
| 4 | 0.189261207 | 10.1.1.2 | 10.1.1.1 | ICMP | 70 | Time-to-live exceeded (Time to live exceeded in transit) |
| 5 | 89.595431119 | 10.1.1.1 | 172.16.1.1 | ICMP | 98 | Echo (ping) request id=0x6250, seq=1/256, ttl=64 (no response found!) |
| 6 | 89.658601894 | 10.1.1.2 | 10.1.1.1 | ICMP | 70 | Destination unreachable (Network unreachable) |
| 7 | 94.834079077 | 30:00:00:00:00:01 | 40:00:00:00:00:03 | ARP | 42 | Who has 10.1.1.2? Tell 10.1.1.1 |
| 8 | 94.924421044 | 40:00:00:00:00:03 | 30:00:00:00:00:01 | ARP | 42 | 10.1.1.2 is at 40:00:00:00:00:03 |

Send a ping from the client to an address that doesn't have a match in the router's forwarding table. So it results in an ICMP destination net unreachable message sent back to the client.

- `client# traceroute -N 1 -n 192.168.100.1`

```
root@njucs-VirtualBox:~/networklab/lab-5-KiraKiraAkira# traceroute -N 1 -n 192.168.100.1
traceroute to 192.168.100.1 (192.168.100.1), 30 hops max, 60 byte packets
 1 10.1.1.2 79.660 ms 106.405 ms 102.045 ms
 2 192.168.100.1 317.071 ms 217.425 ms 220.895 ms
root@njucs-VirtualBox:~/networklab/lab-5-KiraKiraAkira#
```

Wireshark capture from client-eth0. The display filter is empty. The packet list shows twenty-two packets: 13. UDP from 10.1.1.1 to 192.168.100.1, 74 bytes. 14. ICMP from 10.1.1.2 to 10.1.1.1, 70 bytes. 15. UDP from 10.1.1.1 to 192.168.100.1, 74 bytes. 16. ICMP from 192.168.100.1 to 10.1.1.1, 102 bytes. 17. UDP from 10.1.1.1 to 192.168.100.1, 74 bytes. 18. ICMP from 192.168.100.1 to 10.1.1.1, 102 bytes. 19. UDP from 10.1.1.1 to 192.168.100.1, 74 bytes. 20. ICMP from 192.168.100.1 to 10.1.1.1, 102 bytes. 21. ARP from 30.00:00:00:00:01 to 40.00:00:00:00:03, 42 bytes. 22. ARP from 40.00:00:00:00:03 to 30.00:00:00:00:01, 42 bytes. The packet details for packet 16 show 'Destination unreachable (Port unreachable)'.

| No. | Time | Source | Destination | Protocol | Length | Info |
|-----|---------------|-------------------|-------------------|----------|--------|--|
| 13 | 256.144684019 | 10.1.1.1 | 192.168.100.1 | UDP | 74 | 63956 -- 33436 Len=32 |
| 14 | 256.246715493 | 10.1.1.2 | 10.1.1.1 | ICMP | 70 | Time-to-live exceeded (Time to live exceeded in transit) |
| 15 | 256.246832006 | 10.1.1.1 | 192.168.100.1 | UDP | 74 | 53763 -- 33437 Len=32 |
| 16 | 256.563889696 | 192.168.100.1 | 10.1.1.1 | ICMP | 102 | Destination unreachable (Port unreachable) |
| 17 | 256.564005118 | 10.1.1.1 | 192.168.100.1 | UDP | 74 | 60279 -- 33438 Len=32 |
| 18 | 256.781417784 | 192.168.100.1 | 10.1.1.1 | ICMP | 102 | Destination unreachable (Port unreachable) |
| 19 | 256.781536549 | 10.1.1.1 | 192.168.100.1 | UDP | 74 | 46886 -- 33439 Len=32 |
| 20 | 257.002419094 | 192.168.100.1 | 10.1.1.1 | ICMP | 102 | Destination unreachable (Port unreachable) |
| 21 | 250.987285656 | 30:00:00:00:00:01 | 40:00:00:00:00:03 | ARP | 42 | Who has 10.1.1.2? Tell 10.1.1.1 |
| 22 | 261.019436775 | 40:00:00:00:00:03 | 30:00:00:00:00:01 | ARP | 42 | 10.1.1.2 is at 40:00:00:00:00:03 |

| Capturing from router-eth2 | | | | | | |
|--|---------------|-------------------|-------------------|----------|--------|--|
| File Edit View Go Capture Analyze Statistics Telephony Wireless Tools Help | | | | | | |
| Apply a display filter ... <Ctrl-/> Expression... | | | | | | |
| No. | Time | Source | Destination | Protocol | Length | Info |
| 13 | 256.144676035 | 10.1.1.1 | 192.168.100.1 | UDP | 74 | 63956 → 33436 Len=32 |
| 14 | 256.246701151 | 10.1.1.2 | 10.1.1.1 | ICMP | 70 | Time-to-live exceeded (Time to live exceeded in transit) |
| 15 | 256.246823588 | 10.1.1.1 | 192.168.100.1 | UDP | 74 | 53763 → 33437 Len=32 |
| 16 | 256.563876095 | 192.168.100.1 | 10.1.1.1 | ICMP | 102 | Destination unreachable (Port unreachable) |
| 17 | 256.563996783 | 10.1.1.1 | 192.168.100.1 | UDP | 74 | 60279 → 33438 Len=32 |
| 18 | 256.781403881 | 192.168.100.1 | 10.1.1.1 | ICMP | 102 | Destination unreachable (Port unreachable) |
| 19 | 256.781527942 | 10.1.1.1 | 192.168.100.1 | UDP | 74 | 48886 → 33439 Len=32 |
| 20 | 257.002405479 | 192.168.100.1 | 10.1.1.1 | ICMP | 102 | Destination unreachable (Port unreachable) |
| 21 | 258.987281909 | 30:00:00:00:00:01 | 40:00:00:00:00:03 | ARP | 42 | Who has 10.1.1.2? Tell 10.1.1.1 |
| 22 | 261.019421848 | 40:00:00:00:00:03 | 30:00:00:00:00:01 | ARP | 42 | 10.1.1.2 is at 40:00:00:00:00:03 |

4. 总结与感想

Lab 5 is a sum up of lab3 and lab4. In lab3 and lab4, we have modules and in lab5 we need to combine them to deal with different kinds of coming packets. It is much harder than labs before, and I learnt a lot from building a router.

Luckily, I passed all tests, but my router is not that efficient. I simply use a queue to store all waiting packets and only process the head of queue. It is easy to implement, but I doubt its efficiency. Since time is limited, I didn't improve my version. Hope to see a much faster way to handle the waiting the packets.